
Malware classification through Spark

Project No.2 lightning talk
Group: ChickenBurger

Mojtaba Sedigh-Fazli; BahaaEddin AlAila; Shawn(Shengming) Zhang

Introduce to this project:

- 1) rooted from the Microsoft Malware Classification Challenge on Kaggle 2015
- 2) The main task is to classify around 9000 malwares into 9 classes
- 3) The big difficulties are laid on the large size of the malwares, 200GB
- 4) Scala spark code has written to solve this project through AMAZON EC2 service
- 5) Sbt is used for build and package our Scala codes.
- 6) We hooked up the EC2 manually since Flintrock has security group issues

Outlines:

- 1) Data Loading
- 2) Features Engineering
- 3) Training on Random Forest, Gradient Boosted trees and SVMs
- 4) Testing
- 5) Potential improvement and lesson learned

Data Loading

Nasty bug:

```
16/09/21 09:35:17 WARN TaskSetManager: Lost task 0.0 in stage 7.2 (TID 27, 172.19.21.188):  
org.apache.spark.shuffle.FetchFailedException: Too large frame: 7069349343
```

Yo Mama's cell number

Data too large bug → **Remote Shuffle Blocks cannot be larger than 2 GB**

Details: <https://issues.apache.org/jira/browse/SPARK-5928>

Only triggered for remote fetches, not local mode, took us a week to pinpoint

(Thx, lazy execution!)

Once you know the problem, few changes to fix it :

Data Loading -cont'd

Solution

```
val training_set:RDD[(String,Iterable[(String,Double)])] = filesWithPaths.grouped(limit).map(  
    filesAtaTime => {  
        sc.wholeTextFiles(filesAtaTime.mkString(",")) // path/to/file1.asm,path/to/file2.asm,...,path/to/fileN.asm  
        //(path,content)  
        .map({ case (path, content) => (path.split("/").last.split("\\.")(0), processingFunc(content))}).persist(StorageLevel.MEMORY_AND_DISK)  
    })  
}).reduce((a:RDD[(String,Iterable[(String,Double)])], b:RDD[(String,Iterable[(String,Double)])]) => a union b)
```

Load and process 800 files at a time (Matrix Sketching - Liberty), union all

Feature Engineering:

- ASM files:
 - the segCount: which are the # lines of each section (e.g.: text, idata, rdata, data)
 - opcodes Ngrams: which are the 1gram, 2gram, 3gram and 4 grams of the opcode (e.g.: "push-mov"-> 45) (97% accuracy)
- Bytes files:
 - hexadecimal Ngrams: which are the 1gram, 2gram, 3gram and 4 grams of the bytestream
 - Fit all 4 grams into one long, lossless but still gentler on reducebykey, eg:
 - 0x1000000000000000AC
 - 0x2000000000000000AC
 - 0x4000032100000000AC
 - 100 for "??"
 - File size : # of bytes (feature 0x0000000000000000)

Feature Engineering - cont'd

- Discard features appearing less than 5 times per file
- Discard features appearing less than 20 times in all files
- vectorize, .compressed
- Extract once, serialize, refrigerate
- Load from disk and deserialize

Training on Random Forest and Gradient Boosted Trees

1) Random Forest:

Parameter sweep on, 5 fold cross validation. (900 combinations)

Our best 98.2% (sadly did not save the model). And couldn't reproduce because RF randomly sample features.

2) Gradient Boosted Trees and LSVM's:

Spark's mllib, One vs Rest, with RF meta stacker:

Did not give good results(distribution skewed) =>

implemented resampling to equalize the class distribution.

Took forever way past the deadline, and ran out of memory.

(should've resampled hashes only, then joined)

Potential improvement and lesson learned

- 1) Only retain features that showed good class correlation
 - a) Look into RF scores for features
- 2) Extract structure from opcodes (model for loops ...etc), regardless of n-grams. (LSTM ?)
- 3) Write the hex file into binary get McAfee to detect it :p

Questions? and thank you!