

On Building Efficient Recommender System

tooYoungTooSimple

Advisor: Dr. Shannon Quinn

Department of Statistics
The University of Georgia

Overview

Introduction

Feature Engineering

Sample Creation

Statistical Modeling

Ongoing Works

Introduction : Motivation

- **Santander** is challenging us to predict which products their existing customers will use/purchase in the next month based on their past behavior and that of similar customers
- Current recommender system is unsatisfactory with uneven customer experience
- Ongoing **Kaggle** competition

Introduction : Motivation

- Extremely common, and utilized in a variety of areas
- Statistical modeling is a key solution
- Balance between recommendation and customer experience (If recommender system is too aggressive, more products will be recommended correctly, but many “irrelevant” customers will be annoyed.)

Introduction : Data

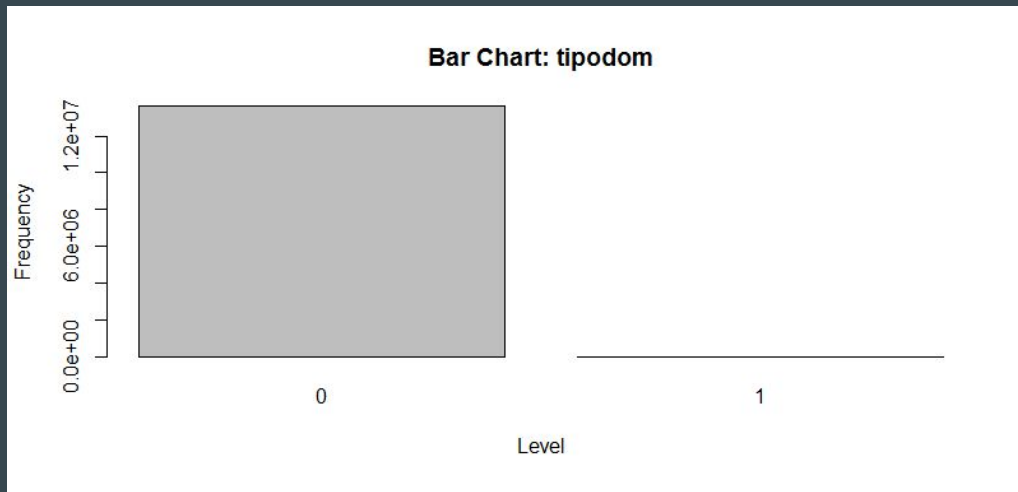
- Observations: 13,647,309 (training) and 929,615 (testing)
- Response: 24 Product Variables
 - Binary and Sparse*
- Predictor: 24 Demographic and Behaviour Variables
- All observations are from 17 months
 - Training data from *2015/01 - 2016/05*
 - Testing data from *2016/06*
- Good News! No new customers appear in testing data

Introduction : Challenging

- Big and Dirty Data
- Imbalance: Very Low Proportion of Products used
- Multiple Products used by Single Customer
- Status of using May Change in Different Patterns
- Missing Values in Response and Predictors

Feature Engineering

- Remove Extremely Imbalanced Variable
 - E.g., AddressType, with 1 category



Feature Engineering

- Impute Missing Values with Different Strategies
 - Predictors: Mode for Categorical Variables
Median for Continuous Variables

E.g. ind_employed

Level	(Missing)	A	B	F	N	S
Frequency	27734	2492	3566	2523	13610977	17



Level	A	B	F	N	S
Frequency	2492	3566	2523	13638711	17

Feature Engineering

- Impute Missing Values with Different Strategies
 - Response: estimated sample expectation
E.g. Ind_nomina_ult1 (Y_22)

Assume a *Bernoulli Distribution*

$$\Pr(Y_{22} = 0) = 0.95$$

$$\Pr(Y_{22} = 1) = 0.05$$

$$E(Y_{22}) = 0.05 * 1 + 0.95 * 0 = 0.05$$


Thus we impute *NAs* in Y_{22} as 0.05

Feature Engineering

- Collapse Levels with Low Frequencies

E.g. canal_entrada

163 levels: 004 007 013 025 K00 KAA ... KAH KAI RED



Feature Engineering

- Collapse Levels with Low Frequencies
E.g. canal_entrada

163 levels: 004 007 013 025 K00 KAA ... KAH KAI RED



13 levels: KHE KAT KFC KHQ KFA KHK KHM KHD
KHN KAS KAG RED and Others

Feature Engineering

- Create 24 Customer Historical Product (*CHP*) Features
 - Past purchase behaviour influence the future purchase
 - Define CHP as weighted sum of past three years' purchase
 - $CHP(Year_i) = \begin{aligned} &(\frac{1}{6}) * Y(Year_i) \\ &+ (\frac{1}{3}) * Y(Year_i - 1) \\ &+ (\frac{1}{2}) * Y(Year_i - 2) \end{aligned}$

If $i = 3$, we use $\frac{1}{3}$ and $\frac{2}{3}$ as two weights

If $i = 2$, we use the previous year's purchase

If $i = 1$, we keep current year's purchase

Sample Creation

- Different response variable has different $\Pr(Y=1)$
- Thus we create 24 samples based on corresponding $\Pr(Y=1)$

For Y_i , keep all observations for $Y_i = 1$

Sampling from the observations with $Y_i = 0$

Such that $\Pr(Y_i = 1) : \Pr(Y_i = 0) = 1:5$

E.g., ind_nom_pens_utm1 (Y_{23})

Original Data: $\#(Y_{23}=1) = 810,085$, $\#(Y_{23}=0) = 12,821,161$

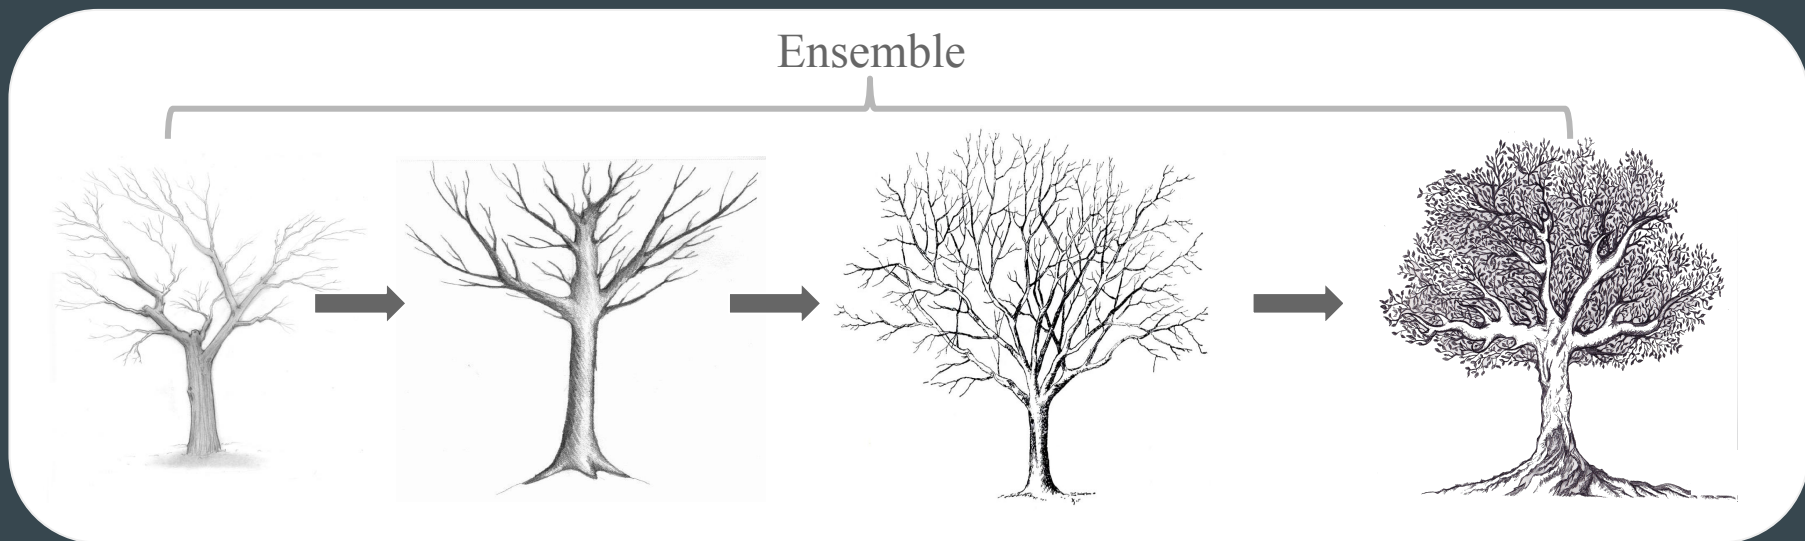
Sample Data: $\#(Y_{23}=1) = 810,085$, $\#(Y_{23}=0) = 4,050,425$

Modeling : Strategy and Procedure

- We have 24 samples in total
- For *sample_i*, taking *Y_i* as Response Variable
taking all *X* and *CHP(-i)* as Predictors
- *model_i* trained with *sample_i* to predict *Y_i* in testing set

Modeling : Gradient Boosted Tree

- Building up the final model by utilizing the error of previous ones.
- Apply l_1 regularization to control overfitting.



Modeling : XGBoost

XGBoost is an optimized distributed gradient boosting library designed to be highly *efficient*, *flexible* and *portable*.

System	exact greedy	approximate global	approximate local	out-of-core	sparsity aware	parallel
XGBoost	yes	yes	yes	yes	yes	yes
SparkMLLib	no	yes	no	no	partially	yes
scikit-learn	yes	no	no	no	no	no
R-GBM	yes	no	no	no	partially	no

Modeling : Cross Validation




5-fold CV for each sample:

- Training model with train set.
- Testing model with test set.
- Aggregate over 5 validations.

Test:



Training:



Ongoing Works

- Will train 24 models with corresponding samples and predict all 24 Responses in testing data
- Kaggle Submission!

Thank you !